

## 14. SQL queries

Principles of Data Science with R

---

Dr. Uma Ravat

PSTAT 10

# Announcement


1. Quiz 6 this week on Canvas
2. HW will be released tomorrow as usual, due next Wednesday

## Next we will see. . .


- Integrity constraints (Primary and Foreign Keys)
- More SQL queries

# SQLite RDBMS and Rstudio

## Open Source Databases



# ChinookDB Entity-Relationship Diagram (ER Diagram)



# Connecting to a DB

- (Install) and load the DBI (Database interface) package
- (Install) and load DBI compliant DataBase Connectivity driver package for the database you will be using
  - For SQLite RDBMS, we will use the SQLite() driver from the RSQLite R package
  - For Postgres RDBMS, use the Postgres() driver from the RPostgres package
  - ...

```
library(DBI)
library(RSQLite)
drv = dbDriver("SQLite") # the driver for the db you want to connect to
chinook_db = dbConnect(drv, # the driver to use
                       dbname="./data/Chinook_Sqlite.sqlite") #path to the db file
```

chinook\_db is an R object that represents a connection to the database file Chinook\_Sqlite.sqlite

# Field metadata

Unlike a data frame, there is extra information in a database table that expresses relational information between tables.

```
dbGetQuery(chinook_db, "pragma table_info(Customer)")
```

##	cid	name	type	notnull	dflt_value	pk
## 1	0	CustomerId	INTEGER	1	NA	1
## 2	1	FirstName	NVARCHAR(40)	1	NA	0
## 3	2	LastName	NVARCHAR(20)	1	NA	0
## 4	3	Company	NVARCHAR(80)	0	NA	0
## 5	4	Address	NVARCHAR(70)	0	NA	0
## 6	5	City	NVARCHAR(40)	0	NA	0
## 7	6	State	NVARCHAR(40)	0	NA	0
## 8	7	Country	NVARCHAR(40)	0	NA	0
## 9	8	PostalCode	NVARCHAR(10)	0	NA	0
## 10	9	Phone	NVARCHAR(24)	0	NA	0
## 11	10	Fax	NVARCHAR(24)	0	NA	0
## 12	11	Email	NVARCHAR(60)	1	NA	0
## 13	12	SupportRepId	INTEGER	0	NA	0

# Primary key

The **primary key** is a *unique identifier* of the rows in a table.

- Two rows cannot have the same primary key

*# paste() is used so that the SQL statement is easy to read*

```
dbGetQuery(chinook_db,
  paste(
    "SELECT CustomerId, FirstName, LastName, City, Country",
    "FROM Customer",
    "LIMIT 2"
  )
)
```

```
## CustomerId FirstName LastName City Country
## 1          1      Luís Gonçalves São José dos Campos Brazil
## 2          2    Leonie Köhler Stuttgart Germany
```

*# No need for paste though*

```
dbGetQuery(chinook_db,
  "SELECT CustomerId, FirstName, LastName, City, Country
  FROM Customer
  LIMIT 2")
```

```
## CustomerId FirstName LastName City Country
## 1          1      Luís Gonçalves São José dos Campos Brazil
## 2          2    Leonie Köhler Stuttgart Germany
```



```
dbExecute(chinook_db,  
  paste("INSERT into Customer",  
        "(CustomerId, FirstName, LastName, Email)",  
        "VALUES",  
        "(1, 'Luis','Armstrong','LuisArmstrong@pstat.ucsb.edu')"  
  )  
)
```

```
## Error: UNIQUE constraint failed: Customer.CustomerId
```

CustomerId is the **primary key** and must be unique.

# Multi-column primary key

Primary key's can consist of multiple columns if it takes multiple columns to identify a row in a table. But, two rows cannot have the same primary key.

*# Single column primary key*

```
dbGetQuery(chinook_db, "pragma table_info(Customer)")
```

##	cid	name	type	notnull	dflt_value	pk
## 1	0	CustomerId	INTEGER	1	NA	1
## 2	1	FirstName	NVARCHAR(40)	1	NA	0
## 3	2	LastName	NVARCHAR(20)	1	NA	0
## 4	3	Company	NVARCHAR(80)	0	NA	0
## 5	4	Address	NVARCHAR(70)	0	NA	0
## 6	5	City	NVARCHAR(40)	0	NA	0
## 7	6	State	NVARCHAR(40)	0	NA	0
## 8	7	Country	NVARCHAR(40)	0	NA	0
## 9	8	PostalCode	NVARCHAR(10)	0	NA	0
## 10	9	Phone	NVARCHAR(24)	0	NA	0
## 11	10	Fax	NVARCHAR(24)	0	NA	0
## 12	11	Email	NVARCHAR(60)	1	NA	0
## 13	12	SupportRepId	INTEGER	0	NA	0


*# Multi column primary key*

```
dbGetQuery(chinook_db, "pragma table_info(PlayListTrack)")
```

##	cid	name	type	notnull	dflt_value	pk
## 1	0	PlaylistId	INTEGER	1	NA	1
## 2	1	TrackId	INTEGER	1	NA	2

# Primary key

Tables are not required to have a primary key, but most do. All the tables in Chinook have a primary key.



# Foreign keys

The relationship between tables is expressed by primary keys and **foreign keys**.

Remember we are working with a relational database, following a relational data model.

```
dbGetQuery(chinook_db,  
           "pragma foreign_key_list(Customer)")
```

```
##   id seq    table      from      to on_update on_delete match  
## 1  0    0 Employee SupportRepId EmployeeId NO ACTION NO ACTION  NONE
```



A foreign key field *points to* the primary key of another table.

# Foreign keys

Foreign keys must either point to an existing value or be NULL.

## To enforce Foreign key constraints in SQLite RDBMS

*# Required for foreign-key support otherwise foreign keys are not enforced*

```
dbExecute(chinook_db, "pragma foreign_keys = on")
```

```
dbGetQuery(chinook_db,  
            "SELECT max(EmployeeId) FROM Employee")
```

```
##      max(EmployeeId)
```

```
## 1                8
```

```
dbGetQuery(chinook_db,  
            "SELECT max(CustomerId) FROM Customer")
```

```
##      max(CustomerId)
```

```
## 1                59
```

```
dbExecute(chinook_db,  
            "INSERT INTO Customer  
            (CustomerId, FirstName, LastName, Email, SupportRepId)  
            VALUES  
            (59, 'Luis', 'Armstrong', 'luisArmstrong@pstat.ucsb.edu', 88)")
```

```
## Error: UNIQUE constraint failed: Customer.CustomerId
```

```
dbExecute(chinook_db,  
  "INSERT INTO Customer  
  (CustomerId, FirstName, LastName, Email, SupportRepId)  
  VALUES  
  (60, 'Luis', 'Armstrong', 'luisArmstrong@pstat.ucsb.edu', 10)")
```

```
## Error: FOREIGN KEY constraint failed
```

# Interpretation of foreign key



- Each customer in Customer table *can* be assigned a support representative
- The support rep is an employee at the store and therefore has a unique id, EmployeeId
- This unique id, EmployeeId, is the primary key of the employee table

Thus real-world relationship is encoded by the relational model using primary and foreign key relationships.



# Integrity Constraints

We have seen two examples of *integrity constraints*:

- Primary keys must be unique (and not NULL)
- Foreign keys must reference existing primary keys or be NULL

These constraints enforce the *integrity* of a database: no bad data or corrupted relationships.

**Keys help maintain the integrity of the data**

# Database Schema

The **schema** of a database describes its *structure*:

- Names of all the tables
- Names of all fields in each table
- Primary key/foreign key relationships between tables
- Other metadata (data types of each field in each table, ...)

Basically everything other than the actual data itself.

Represented via E-R diagrams (Entity relationship)

We have been looking at parts of the schema with the `pragma` keyword.

```
dbGetQuery(chinook_db, "pragma table_info(customer)")
```

## More SQL queries

---

# SELECT

```
dbGetQuery(chinook_db,  
            "SELECT count(*) FROM track")
```

```
##   count(*)  
## 1      3503
```

## What are all the fields for every track?

```
dbListFields(chinook_db, "track")
```

```
## [1] "TrackId" "Name" "AlbumId" "MediaTypeId" "GenreId"  
## [6] "Composer" "Milliseconds" "Bytes" "UnitPrice"
```

```
track_sel <- dbGetQuery(chinook_db,  
                        "SELECT * FROM track")
```

```
str(track_sel)
```

```
## 'data.frame': 3503 obs. of 9 variables:  
## $ TrackId : int 1 2 3 4 5 6 7 8 9 10 ...  
## $ Name : chr "For Those About To Rock (We Salute You)" "Balls to the Wall"  
"Fast As a Shark" "Restless and Wild" ...  
## $ AlbumId : int 1 2 3 3 3 1 1 1 1 1 ...  
## $ MediaTypeId : int 1 2 2 2 2 1 1 1 1 1 ...  
## $ GenreId : int 1 1 1 1 1 1 1 1 1 1 ...  
## $ Composer : chr "Angus Young, Malcolm Young, Brian Johnson" NA "F. Baltes,  
S. Kaufman, U. Dirkschneider & W. Hoffman" "F. Baltes, R.A. Smith-Diesel, S.  
Kaufman, U. Dirkschneider & W. Hoffman" ...  
## $ Milliseconds: int 343719 342562 230619 252051 375418 205662 233926 210834  
203102 263497 ...  
## $ Bytes : int 11170334 5510424 3990994 4331779 6290521 6713451 7636561  
6852860 6599424 8611245 ...
```

Suppose we only want the first five records for TrackId, Name, AlbumId, Milliseconds, Bytes, UnitPrice from Track table

```
dbGetQuery(chinook_db,  
           "SELECT TrackId, Name, AlbumId, Milliseconds, Bytes, UnitPrice  
           FROM track  
           limit 5")
```

```
## TrackId Name AlbumId Milliseconds Bytes  
## 1 1 For Those About To Rock (We Salute You) 1 343719 11170334  
## 2 2 Balls to the Wall 2 342562 5510424  
## 3 3 Fast As a Shark 3 230619 3990994  
## 4 4 Restless and Wild 3 252051 4331779  
## 5 5 Princess of the Dawn 3 375418 6290521  
## UnitPrice  
## 1 0.99  
## 2 0.99  
## 3 0.99  
## 4 0.99  
## 5 0.99
```

## SELECT, expanded

In the first line of SELECT, we can directly specify computations that we want performed

```
SELECT columns or computations  
FROM table  
WHERE condition  
GROUP BY columns  
HAVING condition  
ORDER BY column [ASC | DESC]  
LIMIT offset, count;
```

Main tools for computations:

MIN, MAX, COUNT, SUM, AVG or any math formula

## Example

To calculate the average Milliseconds, Bytes and Max UnitPrice

```
dbGetQuery(chinook_db,  
  paste("SELECT AVG(Milliseconds),  
        AVG(Bytes),MAX(UnitPrice)",  
        "FROM Track"))
```

```
##  AVG(Milliseconds) AVG(Bytes) MAX(UnitPrice)  
##  1          393599.2   33510207          1.99
```

To replicate this simple command on an imported data frame:

```
mean(track_sel$Milliseconds, na.rm=TRUE)
```

```
## [1] 393599.2
```

```
mean(track_sel$Bytes, na.rm=TRUE)
```

```
## [1] 33510207
```

```
max(track_sel$UnitPrice, na.rm=TRUE)
```

```
## [1] 1.99
```

## GROUP BY

We can use the GROUP BY option in SELECT to define aggregation groups

```
dbGetQuery(chinook_db, paste("SELECT AlbumId, AVG(Bytes)",  
                             "FROM Track",  
                             "GROUP BY AlbumId",  
                             "ORDER BY AVG(Bytes) DESC",  
                             "LIMIT 10"))
```

```
##      AlbumId  AVG(Bytes)  
## 1         253  536359244  
## 2         229  535292434  
## 3         227  529469291  
## 4         231  514373372  
## 5         228  512231374  
## 6         254  492670102  
## 7         226  490750393  
## 8         261  453454450  
## 9         251  306109250  
## 10        249  268393262
```

(Note: the order of commands here matters; try switching the order of GROUP BY and ORDER BY, you'll get an error)



We can use AS in the first line of SELECT to rename computed columns

```
dbGetQuery(chinook_db,  
            paste("SELECT AlbumId, AVG(Bytes) AS AvgBytes",  
                  "FROM Track",  
                  "GROUP BY AlbumId",  
                  "ORDER BY AVG(Bytes) DESC",  
                  "LIMIT 10"))
```

	AlbumId	AvgBytes
## 1	253	536359244
## 2	229	535292434
## 3	227	529469291
## 4	231	514373372
## 5	228	512231374
## 6	254	492670102
## 7	226	490750393
## 8	261	453454450
## 9	251	306109250
## 10	249	268393262

## WHERE

```
dbGetQuery(chinook_db,  
  paste("SELECT AlbumId, Avg(Bytes)",  
        "FROM Track",  
        "WHERE AlbumId = 50"  
  ))
```

```
## AlbumId Avg(Bytes)  
## 1      50    30444082
```

We can use the WHERE option in SELECT to specify a subset of the rows to use (*pre-aggregation/pre-calculation*)

```
dbGetQuery(chinook_db,  
  paste("SELECT AlbumId, MediaTypeId,AVG(Bytes) as AvgBytes",  
        "FROM Track",  
        "WHERE AlbumId <= 160",  
        "GROUP BY AlbumId",  
        "ORDER BY AvgBytes DESC",  
        "LIMIT 10"))
```

```
## AlbumId MediaTypeId AvgBytes  
## 1      50          1 30444082  
## 2     138          1 24822832  
## 3     137          1 19120969  
## 4      43          1 16221538  
## 5      97          1 16089011  
## 6     114          1 15975057  
## 7     109          1 15934275  
## 8     113          1 15521017  
## 9     127          1 15194926
```

## HAVING

We can use the HAVING option in SELECT to specify a subset of the rows to display (*post-aggregation/post-calculation*)

```
dbGetQuery(chinook_db,  
  paste("SELECT AlbumId, MediaTypeId,AVG(Bytes) as AvgBytes",  
        "FROM Track",  
        "WHERE AlbumId >= 160",  
        "GROUP BY AlbumId",  
        "HAVING AvgBytes >= 25000000",  
        "ORDER BY AVG(Bytes) DESC",  
        "LIMIT 10"))
```

##	AlbumId	MediaTypeId	AvgBytes
## 1	253	3	536359244
## 2	229	3	535292434
## 3	227	3	529469291
## 4	231	3	514373372
## 5	228	3	512231374
## 6	254	3	492670102
## 7	226	3	490750393
## 8	261	3	453454450
## 9	251	3	306109250
## 10	249	3	268393262

# Disconnecting from the database

After the end of a session, it is good practice to explicitly close your connection.

```
dbDisconnect(chinook_db)

# Try selecting data
dbGetQuery(chinook_db,
            "select CustomerId, FirstName, LastName from Customer")

## Error: Invalid or closed connection
```

Does this remove the database connection `chinook_db` in the R session?

## We saw

- Integrity constraints (Primary and Foreign Keys)
- All parts of a SQL query

SELECT columns or computations

FROM table

WHERE condition

GROUP BY columns

HAVING condition

ORDER BY column [ASC | DESC]

LIMIT offset, count;

- Database tools for R
  - the R packages RSQLite, DBI
  - the database Chinook\_Sqlite.sqlite